

```

import os
import joblib
import json
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.metrics import (
    accuracy_score, classification_report, confusion_matrix
)

MODEL_DIR = "../src/models"
DATA_PATH = "../data/processed/cleaned_combined_data_utf8.csv"

df = pd.read_csv(DATA_PATH)

print(f"Dataset shape: {df.shape}")
print("Categories:", df['category'].value_counts())

X = df['combined_text']
y = df['category']

Dataset shape: (12471, 7)
Categories: category
anxiety      1987
psychosis    1841
addiction     1740
ptsd          994
bipolar       973
bpd           967
ocd           962
autism        902
suicide       827
depression    726
adhd          552
Name: count, dtype: int64

log_reg = joblib.load(os.path.join(MODEL_DIR,
    "baseline_logistic_regression.pkl"))
rand_forest = joblib.load(os.path.join(MODEL_DIR,
    "baseline_random_forest.pkl"))
vectorizer = joblib.load(os.path.join(MODEL_DIR,
    "tfidf_vectorizer.pkl"))

with open(os.path.join(MODEL_DIR, "baseline_results.json"), "r") as f:
    results_summary = json.load(f)

print("Saved training results:")
print(json.dumps(results_summary, indent=2))

```

Saved training results:

```
{
  "timestamp": "20250907_094642",
  "logistic_regression_accuracy": 0.7326653306613227,
  "random_forest_accuracy": 0.6949899799599198,
  "models_saved": [
    "baseline_logistic_regression.pkl",
    "baseline_random_forest.pkl"
  ],
  "vectorizer_saved": "tfidf_vectorizer.pkl"
}
```

```
from sklearn.model_selection import train_test_split
```

```
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)
```

```
X_test_tfidf = vectorizer.transform(X_test)
```

```
from sklearn.metrics import precision_score, recall_score, f1_score
```

```
lr_pred = log_reg.predict(X_test_tfidf)
```

```
print("Logistic Regression Report")
```

```
print(classification_report(y_test, lr_pred))
```

```
lr_acc = accuracy_score(y_test, lr_pred)
```

```
lr_prec_macro = precision_score(y_test, lr_pred, average="macro")
```

```
lr_recall_macro = recall_score(y_test, lr_pred, average="macro")
```

```
lr_f1_macro = f1_score(y_test, lr_pred, average="macro")
```

```
lr_prec_weighted = precision_score(y_test, lr_pred,
average="weighted")
```

```
lr_recall_weighted = recall_score(y_test, lr_pred, average="weighted")
```

```
lr_f1_weighted = f1_score(y_test, lr_pred, average="weighted")
```

```
print(f"\nLogistic Regression Metrics:")
```

```
print(f"Accuracy: {lr_acc:.4f}")
```

```
print(f"Precision (Macro): {lr_prec_macro:.4f}")
```

```
print(f"Recall (Macro): {lr_recall_macro:.4f}")
```

```
print(f"F1 Score (Macro): {lr_f1_macro:.4f}")
```

```
print(f"Precision (Weighted): {lr_prec_weighted:.4f}")
```

```
print(f"Recall (Weighted): {lr_recall_weighted:.4f}")
```

```
print(f"F1 Score (Weighted): {lr_f1_weighted:.4f}")
```

Logistic Regression Report

	precision	recall	f1-score	support
addition	0.90	0.90	0.90	348
adhd	0.68	0.69	0.68	110
anxiety	0.78	0.75	0.77	398

autism	0.59	0.72	0.65	180
bipolar	0.71	0.64	0.67	195
bpd	0.74	0.77	0.76	194
depression	0.39	0.46	0.42	145
ocd	0.83	0.82	0.83	193
psychosis	0.85	0.71	0.77	368
ptsd	0.75	0.72	0.74	199
suicide	0.54	0.66	0.59	165
accuracy			0.73	2495
macro avg	0.71	0.71	0.71	2495
weighted avg	0.74	0.73	0.74	2495

Logistic Regression Metrics:

Accuracy: 0.7327

Precision (Macro): 0.7065

Recall (Macro): 0.7120

F1 Score (Macro): 0.7070

Precision (Weighted): 0.7443

Recall (Weighted): 0.7327

F1 Score (Weighted): 0.7362

```
rf_pred = rand_forest.predict(X_test_tfidf)
```

```
print("Random Forest Report")
```

```
print(classification_report(y_test, rf_pred))
```

```
rf_acc = accuracy_score(y_test, rf_pred)
```

```
rf_prec_macro = precision_score(y_test, rf_pred, average="macro")
```

```
rf_recall_macro = recall_score(y_test, rf_pred, average="macro")
```

```
rf_f1_macro = f1_score(y_test, rf_pred, average="macro")
```

```
rf_prec_weighted = precision_score(y_test, rf_pred,
average="weighted")
```

```
rf_recall_weighted = recall_score(y_test, rf_pred, average="weighted")
```

```
rf_f1_weighted = f1_score(y_test, rf_pred, average="weighted")
```

```
print(f"\nRandom Forest Metrics:")
```

```
print(f"Accuracy: {rf_acc:.4f}")
```

```
print(f"Precision (Macro): {rf_prec_macro:.4f}")
```

```
print(f"Recall (Macro): {rf_recall_macro:.4f}")
```

```
print(f"F1 Score (Macro): {rf_f1_macro:.4f}")
```

```
print(f"Precision (Weighted): {rf_prec_weighted:.4f}")
```

```
print(f"Recall (Weighted): {rf_recall_weighted:.4f}")
```

```
print(f"F1 Score (Weighted): {rf_f1_weighted:.4f}")
```

Random Forest Report

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

addiction	0.82	0.87	0.84	348
adhd	0.69	0.67	0.68	110
anxiety	0.81	0.72	0.76	398
autism	0.41	0.76	0.53	180
bipolar	0.69	0.63	0.66	195
bpd	0.75	0.69	0.72	194
depression	0.40	0.38	0.39	145
ocd	0.79	0.80	0.80	193
psychosis	0.90	0.59	0.71	368
ptsd	0.81	0.70	0.75	199
suicide	0.48	0.68	0.56	165
accuracy			0.69	2495
macro avg	0.69	0.68	0.67	2495
weighted avg	0.73	0.69	0.70	2495

Random Forest Metrics:

Accuracy: 0.6950

Precision (Macro): 0.6869

Recall (Macro): 0.6805

F1 Score (Macro): 0.6736

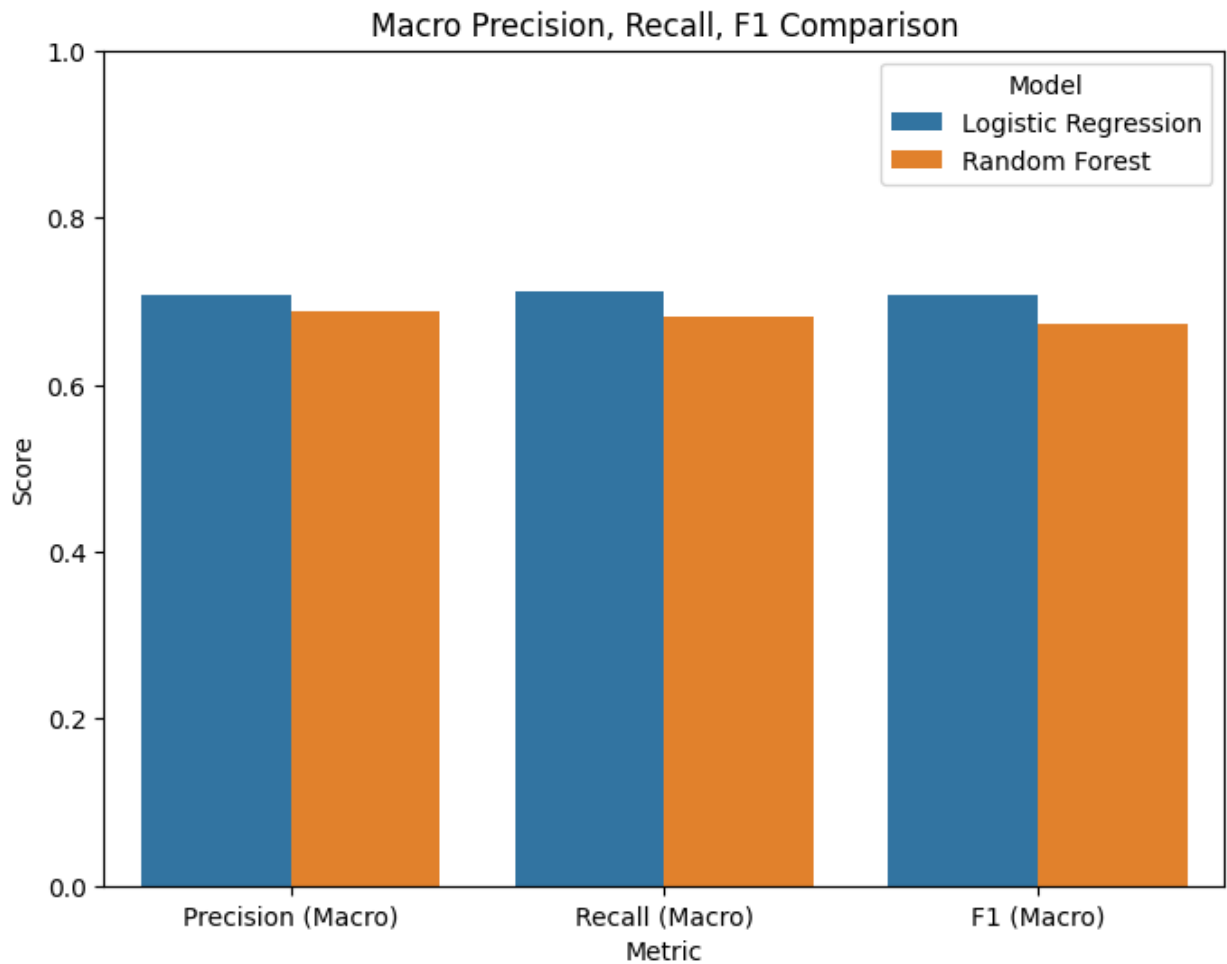
Precision (Weighted): 0.7301

Recall (Weighted): 0.6950

F1 Score (Weighted): 0.7019

```
metrics_df = pd.DataFrame({
    "Model": ["Logistic Regression", "Logistic Regression", "Logistic Regression",
              "Random Forest", "Random Forest", "Random Forest"],
    "Metric": ["Precision (Macro)", "Recall (Macro)", "F1 (Macro)",
               "Precision (Macro)", "Recall (Macro)", "F1 (Macro)"],
    "Score": [lr_prec_macro, lr_recall_macro, lr_f1_macro,
               rf_prec_macro, rf_recall_macro, rf_f1_macro]
})
```

```
plt.figure(figsize=(8,6))
sns.barplot(x="Metric", y="Score", hue="Model", data=metrics_df)
plt.ylim(0,1)
plt.title("Macro Precision, Recall, F1 Comparison")
plt.show()
```



```
sample_df = pd.DataFrame({
    "Text": X_test.sample(5, random_state=42)
})

sample_df["Actual"] = y_test.loc[sample_df.index]
sample_df["LogReg_Pred"] =
log_reg.predict(vectorizer.transform(sample_df["Text"]))
sample_df["RF_Pred"] =
rand_forest.predict(vectorizer.transform(sample_df["Text"]))
```

sample_df

	Text	Actual
LogReg_Pred \		
4093	Tell me what you're currently anxious about an...	anxiety
anxiety		
10910	Feelings of re-experiencing are distressing af...	ptsd
ptsd		
6152	App for us bipolars I want to build an app tha...	bipolar
bipolar		

3930 23 and never had an actual job, feeling lost. ... anxiety
depression
1962 Mallinkrodt Generic = worthless Does anyone el... adhd
adhd

RF_Pred
4093 anxiety
10910 ptsd
6152 bipolar
3930 bpd
1962 adhd